

CS315: DATABASE SYSTEMS INDEXING

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

2nd semester, 2019-20

Tue, Wed 12:00-13:15

- Indexing is used to speed up search
- A **search key** is used
- An **index file** consists of records or **index entries** which has two fields
 - 1 Search key: Attribute that is used for searching
 - 2 Pointer to the entire object or tuple
- Index files should be smaller than data files

- Indexing is used to speed up search
- A **search key** is used
- An **index file** consists of records or **index entries** which has two fields
 - 1 Search key: Attribute that is used for searching
 - 2 Pointer to the entire object or tuple
- Index files should be smaller than data files
- Index evaluation metrics
 - Search time
 - Modification overhead
 - Space overhead

- Indexing is used to speed up search
- A **search key** is used
- An **index file** consists of records or **index entries** which has two fields
 - 1 Search key: Attribute that is used for searching
 - 2 Pointer to the entire object or tuple
- Index files should be smaller than data files
- Index evaluation metrics
 - Search time
 - Modification overhead
 - Space overhead
- Two basic types of indices
 - 1 **Ordered index**: search keys are organized according to some order
 - 2 **Hash index**: search keys are organized according to a hash function

Static Hashing

- A **hash function** maps a key to a bucket
- A **bucket** is a unit of storage
- It is typically a disk block
- A key may need to be searched sequentially inside a bucket
- Results in **hash file organization**
- Example: mod n where n is the number of buckets

Hash Function

- Two important qualities of an ideal hash function
- **Uniform**: Total number of keys from the domain is spread uniformly over all the buckets
- **Random**: Number of keys in each bucket is same irrespective of the actual distribution of keys

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing:**

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**
- **Open addressing**:

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
 - Address at i^{th} attempt for key k is $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
 - **Linear probing**:

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
 - Address at i^{th} attempt for key k is $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
 - **Linear probing**: Interval between buckets is fixed
 - $h(k, i) = (h(k) + c \cdot i) \bmod m$
 - **Quadratic probing**:

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
 - Address at i^{th} attempt for key k is $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
 - **Linear probing**: Interval between buckets is fixed
 - $h(k, i) = (h(k) + c \cdot i) \bmod m$
 - **Quadratic probing**: Interval between buckets increase linearly
 - $h(k, i) = (h(k) + c \cdot i \cdot i) \bmod m$
 - **Double hashing**:

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
 - Address at i^{th} attempt for key k is $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
 - **Linear probing**: Interval between buckets is fixed
 - $h(k, i) = (h(k) + c \cdot i) \bmod m$
 - **Quadratic probing**: Interval between buckets increase linearly
 - $h(k, i) = (h(k) + c \cdot i \cdot i) \bmod m$
 - **Double hashing**: Interval computed using another hash function h'
 - $h(k, i) = (h(k) + h'(k) \cdot i) \bmod m$

Bucket Overflow

- Buckets may overflow at runtime due to
 - Skew in distribution of keys
 - Non-uniformity of hash function
- Probability can only be reduced, but not eliminated
- Handled using **overflow buckets**
- **Closed addressing**: Overflow buckets are linked together
 - Also called **separate chaining** or **chaining**
- **Open addressing**: Excess keys assigned to some other bucket and total number of buckets is kept fixed
 - Address at i^{th} attempt for key k is $h(k, i) = (h(k) + o_i \cdot i) \bmod m$
 - **Linear probing**: Interval between buckets is fixed
 - $h(k, i) = (h(k) + c \cdot i) \bmod m$
 - **Quadratic probing**: Interval between buckets increase linearly
 - $h(k, i) = (h(k) + c \cdot i \cdot i) \bmod m$
 - **Double hashing**: Interval computed using another hash function h'
 - $h(k, i) = (h(k) + h'(k) \cdot i) \bmod m$
- Changing size of a database is a problem
- Periodic re-hashing is the only solution
- **Dynamic hashing**: h changes dynamically but deterministically

Dynamic Hashing

- Organize overflow buckets as binary trees
- m binary trees for m primary pages
- $h_0(k)$ produces index of primary page
- Particular access structure for binary trees
- Family of functions $g(k) = \{h_1(k), \dots, h_i(k), \dots\}$
- Each $h_i(k)$ produces a bit
- At level i , if $h_i(k) = 0$, take left branch, otherwise right branch
- Example: bit representation

Ordered Index

- **Index sequential file**: ordered sequential file with an index

Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
 - Generally primary key
- **Secondary index** or **non-clustering index**: any other index

Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
 - Generally primary key
- **Secondary index** or **non-clustering index**: any other index
- **Dense index**: index record appears for every key
 - Secondary index must be dense
- **Sparse index**: index record appears for only some keys
 - Records must be sequentially ordered by key
 - To locate k , find largest key $< k$, and then sequential search from there

Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
 - Generally primary key
- **Secondary index** or **non-clustering index**: any other index
- **Dense index**: index record appears for every key
 - Secondary index must be dense
- **Sparse index**: index record appears for only some keys
 - Records must be sequentially ordered by key
 - To locate k , find largest key $< k$, and then sequential search from there
 - Deletion: Replaced by next key
 - Insertion: Inserted only if entries exceed a block
- Sparse index requires less space and less maintenance but more searching time

Ordered Index

- **Index sequential file**: ordered sequential file with an index
- **Primary index** or **clustering index**: index whose search key specifies the sequential order of the file
 - Generally primary key
- **Secondary index** or **non-clustering index**: any other index
- **Dense index**: index record appears for every key
 - Secondary index must be dense
- **Sparse index**: index record appears for only some keys
 - Records must be sequentially ordered by key
 - To locate k , find largest key $< k$, and then sequential search from there
 - Deletion: Replaced by next key
 - Insertion: Inserted only if entries exceed a block
- Sparse index requires less space and less maintenance but more searching time
- **Multilevel index**: primary index does not fit in memory
 - **Outer index**: Sparse primary index
 - **Inner index**: Dense primary index file


- Balanced hierarchical data structure
- Keys (and associated objects) are in secondary storage, i.e., disk
- A B-tree of order Θ has the following properties:
 - 1 Leaf nodes are in same level, i.e., the tree is balanced
 - 2 Root has at least 1 key
 - 3 Other internal nodes have between Θ and 2Θ keys
 - 4 An internal node with k keys have $k + 1$ children
 - 5 Child pointers in leaf nodes are null
- Branching factor is between $\Theta + 1$ and $2\Theta + 1$
- Pointer to the object corresponding to a key is stored alongside

B+-Tree


- Most important variety of B-tree
- Internal nodes do *not* contain pointers to objects
- Often siblings are connected by pointers to avoid parent traversal

B+-Tree

- Most important variety of B-tree
- Internal nodes do *not* contain pointers to objects
- Often siblings are connected by pointers to avoid parent traversal




B-tree node




B+-tree node

B+-Tree

- Most important variety of B-tree
- Internal nodes do *not* contain pointers to objects
- Often siblings are connected by pointers to avoid parent traversal



B-tree node



B+-tree node

- More keys can fit in a B+-tree
- Height may be less

Order

- Order mainly determined by disk page size

Order

- Order mainly determined by disk page size
- Disk page capacity is C bytes
- Key consumes γ bytes
- Pointer to key (or object) requires η bytes

Order

- Order mainly determined by disk page size
- Disk page capacity is C bytes
- Key consumes γ bytes
- Pointer to key (or object) requires η bytes
- For B+-tree

Order

- Order mainly determined by disk page size
- Disk page capacity is C bytes
- Key consumes γ bytes
- Pointer to key (or object) requires η bytes
- For B+-tree

$$C \geq 2.\Theta.\gamma + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B+-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + \eta)} \right\rfloor$$

Order

- Order mainly determined by disk page size
- Disk page capacity is C bytes
- Key consumes γ bytes
- Pointer to key (or object) requires η bytes
- For B+-tree

$$C \geq 2.\Theta.\gamma + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B+-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + \eta)} \right\rfloor$$

- For B-tree

Order

- Order mainly determined by disk page size
- Disk page capacity is C bytes
- Key consumes γ bytes
- Pointer to key (or object) requires η bytes
- For B+-tree

$$C \geq 2.\Theta.\gamma + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B+-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + \eta)} \right\rfloor$$

- For B-tree

$$C \geq 2.\Theta.\gamma + 2.\Theta.\eta + (2.\Theta + 1).\eta$$
$$\Rightarrow \Theta_{\text{B-tree}} = \left\lfloor \frac{C - \eta}{2(\gamma + 2\eta)} \right\rfloor$$

Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?

Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?

- $\Theta_{\text{B+-tree}} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$

Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?

- $\Theta_{B+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$

- $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$

Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?
 - $\Theta_{\text{B+-tree}} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$
 - $\Theta_{\text{B-tree}} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$
- What is the height of the B+-tree and the B-tree for 3×10^7 keys?

Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?
 - $\Theta_{B+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$
 - $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$
- What is the height of the B+-tree and the B-tree for 3×10^7 keys?
 - Height of B+-tree is $\lceil \log_{2 \times 170}(3 \times 10^7) \rceil = 3$

Example

- What is the order of a B+-tree and a B-tree with a page size of 4 KB indexing keys of 8 bytes each, and having pointers of size 4 bytes?
 - $\Theta_{B+-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+4)} \right\rfloor = 170$
 - $\Theta_{B-tree} = \left\lfloor \frac{4 \times 1024 - 4}{2(8+2 \times 4)} \right\rfloor = 127$
- What is the height of the B+-tree and the B-tree for 3×10^7 keys?
 - Height of B+-tree is $\lceil \log_{2 \times 170}(3 \times 10^7) \rceil = 3$
 - Height of B-tree is $\lceil \log_{2 \times 127}(3 \times 10^7) \rceil = 4$

Indexing Multiple Attributes

- Search keys having more than one attribute are called **composite search keys**
- Separate indices may be used
 - Union, intersection, etc. of individual results
- Multi-dimensional indexing
 - Indexing is specified by hyper-rectangles

Indexing Multiple Attributes

- Search keys having more than one attribute are called **composite search keys**
- Separate indices may be used
 - Union, intersection, etc. of individual results
- Multi-dimensional indexing
 - Indexing is specified by hyper-rectangles
 - Space-partitioning
 - **Quadtree, KD-tree, K-d-B-tree**: Extension of BST

Indexing Multiple Attributes

- Search keys having more than one attribute are called **composite search keys**
- Separate indices may be used
 - Union, intersection, etc. of individual results
- Multi-dimensional indexing
 - Indexing is specified by hyper-rectangles
 - Space-partitioning
 - **Quadtree**, **KD-tree**, **K-d-B-tree**: Extension of BST
 - Data-partitioning
 - **R-tree**: Extension of B+-tree
 - Uses **minimum bounding rectangles (MBRs)**

Bitmap Index

- Attribute domain consists of a small number of distinct values
- A **bitmap** or a **bit vector** is an array of bits
- Each distinct value has an array of the size of the number of tuples
 - If the i -th bit is 1, tuple i has that value

Gender	Grade
Male	C
Female	A
Female	C
Male	D
Male	A

- Two sets of bit vectors
 - Male = (10011), Female = (01100)
 - A = (01001), B = (00000), C = (10100), D = (00010)

Bitmap Operations

- Queries are answered using bitmap operations
- Example: Find the male student who got 'D'
 - $\text{Bitmap}(\text{Male}) \text{ AND } \text{Bitmap}(\text{D})$
- Null values require a special bitmap for null
- O/S allows efficient bitmap operations when they are packed in word sizes