

CS315: DATABASE SYSTEMS STRUCTURED QUERY LANGUAGE (SQL)

Arnab Bhattacharya

`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,
Indian Institute of Technology, Kanpur

`http://web.cse.iitk.ac.in/~cs315/`

2nd semester, 2019-20

Tue, Wed 12:00-13:15

Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases

Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
 - Can access and manipulate data stored as a particular data model

Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
 - Can access and manipulate data stored as a particular data model
- *Declarative language*
 - Specifies what to do, but not how to do

Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
 - Can access and manipulate data stored as a particular data model
- *Declarative language*
 - Specifies what to do, but not how to do
- Is also a **data definition language (DDL)**
 - Defines database relations and schemas

Structured Query Language (SQL)

- **SQL** is a *querying* language for relational databases
- Is a **data manipulation language (DML)**
 - Can access and manipulate data stored as a particular data model
- *Declarative language*
 - Specifies what to do, but not how to do
- Is also a **data definition language (DDL)**
 - Defines database relations and schemas
- SQL has evolved widely after its first inception
 - Supports lots of extra operations that are non-standard

Example Schema

- course (code, title, *ctype*, webpage)
- coursetype (ctype, *dept*)
- faculty (fid, name, *dept*, designation)
- department (deptid, name)
- semester (yr, half)
- offering (*coursecode*, yr, half, instructor)
- student (roll, name, *dept*, cpi)
- program (roll, *p*type)
- registration (*coursecode*, roll, yr, half, gradecode)
- grade (gradecode, value)

Creating Relation Schemas

- **create table**: create table

$r(A_1 D_1 C_1, \dots, A_n D_n C_n, (IC_1), \dots, (IC_k))$

- r is the name of the relation
- Each A_i is an attribute name whose data type or domain is specified by D_i
- C_i specifies constraints or settings (if any)
- IC_j represents integrity constraints (if any)

- Example

```
create table faculty (  
    fid integer primary key,  
    name varchar(50) not null,  
    dept integer,  
    designation varchar(3)  
);
```


Data Types in SQL

- *char(n)*: fixed-length character string
- *varchar(n)*: variable-length character string, up to n
- *integer* or *int*: integer
- *smallint*: short integer
- *numeric(n,d)*: floating-point number with a total of n digits of which d is after the decimal point
- *real*: single-precision floating-point number
- *double precision*: double-precision floating-point number
- *float(n)*: floating-point number with at least n digits

Data Types in SQL

- *char(n)*: fixed-length character string
- *varchar(n)*: variable-length character string, up to n
- *integer* or *int*: integer
- *smallint*: short integer
- *numeric(n,d)*: floating-point number with a total of n digits of which d is after the decimal point
- *real*: single-precision floating-point number
- *double precision*: double-precision floating-point number
- *float(n)*: floating-point number with at least n digits
- *date*: yyyy-mm-dd format
- *time*: hh:mm:ss format
- *time(i)*: hh:mm:ss:i . . . i format with additional i digits for fraction of a second
- *timestamp*: both date and time
- *interval*: relative value in either year-month or day-time format

Other Data Types

- User-defined data type

```
create type cpi as numeric(3,1);
```

- Large objects such as images, videos, strings can be stored as
 - **blob**: binary large object
 - **clob**: character large object
 - A pointer to the object is stored in the relation, and not the object itself

Other Data Types

- User-defined data type

```
create type cpi as numeric(3,1);
```

- Large objects such as images, videos, strings can be stored as
 - **blob**: binary large object
 - **clob**: character large object
 - A pointer to the object is stored in the relation, and not the object itself
- User-defined domain

```
create domain name as varchar(50) not null;
```

Constraints

- Can be specified for each attribute as well as separately
 - *not null*: the attribute cannot be null
 - Requires some value while inserting as otherwise null is the default
 - *primary key* (A_i, \dots, A_j): automatically ensures not null
 - *default n*: defaults to n if no value is specified
 - *unique*: specifies that this is a candidate key
 - *foreign key*: specifies as a foreign key and the relation it refers to
 - *check P*: predicate P must be satisfied

```
create table faculty (  
    fid integer ,  
    name varchar(50) not null ,  
    dept integer ,  
    designation varchar(3) default 'AP' ,  
    primary key fid ,  
    foreign key (dept) references department(dept) ,  
    check (fid >= 0)  
);
```

Deleting or Modifying a Relation Schema

- **drop table**: **drop table** *r* deletes the table from the database
 - Must satisfy other constraints already applied

- Example

drop table faculty ;

Deleting or Modifying a Relation Schema

- **drop table**: **drop table** r deletes the table from the database
 - Must satisfy other constraints already applied

- Example

```
drop table faculty;
```

- **alter table**: **alter table** r **add** A D C
 - Adds attribute A with data type D at the end
 - C specifies constraints on A (if any)
 - Must satisfy other constraints already applied
- **alter table**: **alter table** r **drop** A
 - Deletes attribute A from all tuples
 - Must satisfy other constraints already applied

- Example

```
alter table faculty add room varchar(10);  
alter table course drop webpage;
```

Basic Query Structure

- SQL is based on *relational algebra*
- A *basic* SQL query is of the form
 select A_1, \dots, A_n
 from r_1, \dots, r_m
 where P
- Each r_i is a relation
- Each A_j is an attribute from one of r_1, \dots, r_m
- P is a predicate involving attributes and constants
- **where** can be left out, which then means *true*
- Result is a relation with the schema (A_1, \dots, A_n)

Basic Query Structure

- SQL is based on *relational algebra*
- A *basic* SQL query is of the form
 select A_1, \dots, A_n
 from r_1, \dots, r_m
 where P
- Each r_i is a relation
- Each A_j is an attribute from one of r_1, \dots, r_m
- P is a predicate involving attributes and constants
- *where* can be left out, which then means *true*
- Result is a relation with the schema (A_1, \dots, A_n)
- Is equivalent to the relational algebra query

Basic Query Structure

- SQL is based on *relational algebra*
- A *basic* SQL query is of the form
 select A_1, \dots, A_n
 from r_1, \dots, r_m
 where P
- Each r_i is a relation
- Each A_j is an attribute from one of r_1, \dots, r_m
- P is a predicate involving attributes and constants
- **where** can be left out, which then means *true*
- Result is a relation with the schema (A_1, \dots, A_n)
- Is equivalent to the relational algebra query
 $\Pi_{A_1, \dots, A_n}(\sigma_P(r_1 \times \dots \times r_m))$

Multisets

- SQL relations are **multisets** or **bags** of tuples and not sets
- Consequently, there may be two identical tuples
- This is the biggest distinction with relational algebra

Multisets

- SQL relations are **multisets** or **bags** of tuples and not sets
- Consequently, there may be two identical tuples
- This is the biggest distinction with relational algebra
- The set behavior can be enforced by the keyword **unique**
- In a query, keyword **distinct** achieves the same effect
- Opposite is keyword **all**, which is *default*

Select

- Lists attributes in the final output

Select

- Lists attributes in the final output
- Example: Find codes of courses offered in 2018

Select

- Lists attributes in the final output
- Example: Find codes of courses offered in 2018

```
select coursecode  
from offering  
where yr = 2018;
```

- Case-insensitive
- **select *** chooses all attributes
- To eliminate duplicates, use **select distinct ...**
- Otherwise, by default is **select all ...**

Select

- Lists attributes in the final output
- Example: Find codes of courses offered in 2018

```
select coursecode  
from offering  
where yr = 2018;
```

- Case-insensitive
- **select** * chooses all attributes
- To eliminate duplicates, use **select distinct ...**
- Otherwise, by default is **select all ...**
- Can contain arithmetic expressions

```
select coursecode, yr - 1959  
from offering  
where yr = 2018;
```


From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations

From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

```
select title  
from course, offering  
where course.code = offering.coursecode and yr = 2018;
```

From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

```
select title  
from course, offering  
where course.code = offering.coursecode and yr = 2018;
```

- When two relations contain attributes of the same name, qualification is needed to remove ambiguity
- Example: Find name of students in B.Tech. program

From

- Lists relations from where attributes will be listed
- Corresponds to Cartesian product of the relations
- Example: Find title of courses offered in 2018

```
select title  
from course, offering  
where course.code = offering.coursecode and yr = 2018;
```

- When two relations contain attributes of the same name, qualification is needed to remove ambiguity
- Example: Find name of students in B.Tech. program

```
select student.name  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.';
```

Where

- Specifies conditions that the result tuples must satisfy

Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

- May use **and**, **or** and **not** to connect predicates

```
select coursecode  
from offering  
where yr = 2018 and instructor = 10;
```


Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

- May use **and**, **or** and **not** to connect predicates

```
select coursecode  
from offering  
where yr = 2018 and instructor = 10;
```

- Unused clause is equivalent to **where true**

```
select coursecode  
from offering;
```

Where

- Specifies conditions that the result tuples must satisfy
- Example

```
select coursecode  
from offering  
where yr = 2018;
```

- May use **and**, **or** and **not** to connect predicates

```
select coursecode  
from offering  
where yr = 2018 and instructor = 10;
```

- Unused clause is equivalent to **where true**

```
select coursecode  
from offering;
```

- SQL allows **between** operator (includes both)

```
select coursecode  
from offering  
where yr between 2016 and 2018;
```

Rename Operation

- SQL allows renaming of relations and attributes to remove ambiguity
- Keyword **as** is used
- Example

```
select student.roll as rollnumber  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.';
```

Rename Operation

- SQL allows renaming of relations and attributes to remove ambiguity
- Keyword **as** is used
- Example

```
select student.roll as rollnumber  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.' ;
```

- Renaming is necessary when the same relation needs to be used twice
- Example: Find names of students whose cpi is greater than that of "ABC"

Rename Operation

- SQL allows renaming of relations and attributes to remove ambiguity
- Keyword **as** is used
- Example

```
select student.roll as rollnumber  
from student, program  
where student.roll = program.roll and program.ptype =  
    'B.Tech.';
```

- Renaming is necessary when the same relation needs to be used twice
- Example: Find names of students whose cpi is greater than that of "ABC"

```
select T.name  
from student as T, student as S  
where T.cpi > S.cpi and S.name = 'ABC';
```

- **as** can be omitted by simply stating **student T**

String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring

String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```


String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```

- Example: Find departments with the name “?E”

String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```

- Example: Find departments with the name “?E”

```
select *  
from department  
where name like ‘ ‘_E ’ ’;
```

String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘ ‘%Engineering%’ ’;
```

- Example: Find departments with the name “?E”

```
select *  
from department  
where name like ‘ ‘_E’ ’;
```

- Example: Find the department whose name is “_E”

String Operations

- Supports string matching in addition to equality of two strings
- Uses **like** to match patterns specified using special characters
 - **_**: matches any character
 - **%**: matches any substring
- Example: Find all departments having “Engineering” in its name

```
select *  
from department  
where name like ‘“%Engineering%” ’;
```

- Example: Find departments with the name “?E”

```
select *  
from department  
where name like ‘“_E” ’;
```

- Example: Find the department whose name is “_E”

```
select *  
from department  
where name = ‘“_E” ’;
```

Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**

Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**
 - For display purposes only and has no actual effect
- Example: Order departments by name

Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**
 - For display purposes only and has no actual effect
- Example: Order departments by name

```
select *  
from department  
order by name;
```

- Use **desc** to obtain tuples in descending order

```
select *  
from department  
order by name desc;
```

Ordering of Tuples

- Tuples in the final relation can be ordered using **order by**
 - For display purposes only and has no actual effect
- Example: Order departments by name

```
select *  
from department  
order by name;
```

- Use **desc** to obtain tuples in descending order

```
select *  
from department  
order by name desc;
```

- Default is ascending order (**asc**)

Set Operations

- Operators **union**, **intersect** and **except** correspond to \cup , \cap , $-$
- *Eliminates* duplicates
- For multiset operations, i.e., to retain duplicates, use **all** after the operations
- Example: Find names of all faculty members and students

Set Operations

- Operators **union**, **intersect** and **except** correspond to \cup , \cap , $-$
- *Eliminates* duplicates
- For multiset operations, i.e., to retain duplicates, use **all** after the operations
- Example: Find names of all faculty members and students

```
(select name from faculty)  
union  
(select name from student);
```
- Example: Find names of faculty members not found in students

Set Operations

- Operators **union**, **intersect** and **except** correspond to \cup , \cap , $-$
- *Eliminates* duplicates
- For multiset operations, i.e., to retain duplicates, use **all** after the operations
- Example: Find names of all faculty members and students

```
(select name from faculty)  
union  
(select name from student);
```

- Example: Find names of faculty members not found in students

```
(select name from faculty)  
except  
(select name from student);
```

Aggregate Functions

- Five operations that work on multisets: **avg**, **min**, **max**, **sum**, **count**
- Example: Find the average cpi of students

Aggregate Functions

- Five operations that work on multisets: **avg**, **min**, **max**, **sum**, **count**
- Example: Find the average cpi of students

```
select avg(cpi)  
from student;
```

- For set operations, use **distinct**
- Example: Find the total number of students

Aggregate Functions

- Five operations that work on multisets: **avg**, **min**, **max**, **sum**, **count**
- Example: Find the average cpi of students

```
select avg(cpi)  
from student;
```

- For set operations, use **distinct**
- Example: Find the total number of students

```
select count(distinct roll)  
from student;
```

Grouping

- To apply aggregate operations on separate groups, use **group by**
- The aggregate operator is applied on *each* group separately
- Example: Find the number of students in each department

Grouping

- To apply aggregate operations on separate groups, use **group by**
- The aggregate operator is applied on *each* group separately
- Example: Find the number of students in each department

```
select dept, count(distinct roll)  
from student  
group by dept;
```


Grouping

- To apply aggregate operations on separate groups, use **group by**
- The aggregate operator is applied on *each* group separately
- Example: Find the number of students in each department

```
select dept, count(distinct roll)  
from student  
group by dept;
```

- Attributes in **select** clause outside of aggregate functions *must* appear in **group by** list

Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

```
select coursecode, avg(grade)
from registration
group by coursecode
having count(roll) >= 5;
```

- The predicate in **having** is applied *after* forming groups whereas the predicate in **where** is applied *before* doing so

Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

```
select coursecode, avg(grade)
from registration
group by coursecode
having count(roll) >= 5;
```

- The predicate in **having** is applied *after* forming groups whereas the predicate in **where** is applied *before* doing so
- Example: Find average grade in each course of type 4 where number of students is at least 5

Qualifying Groups

- In order to select certain groups, use **having** clause
- Only those groups satisfying **having** clause appear in the result
- Example: Find average grade in each course where number of students is at least 5

```
select coursecode, avg(grade)
from registration
group by coursecode
having count(roll) >= 5;
```

- The predicate in **having** is applied *after* forming groups whereas the predicate in **where** is applied *before* doing so
- Example: Find average grade in each course of type 4 where number of students is at least 5

```
select coursecode, avg(grade)
from registration, course
where registration.coursecode = course.code and ctype = 4
group by coursecode
having count(roll) >= 5;
```

Null

- *null* signifies missing or unknown value
- The predicates *is null* and *is not null* can be used to check for null values
- Example: find courses that do not have a webpage

Null

- *null* signifies missing or unknown value
- The predicates **is null** and **is not null** can be used to check for null values
- Example: find courses that do not have a webpage

```
select code  
from course  
where webpage is null;
```

- *null* signifies missing or unknown value
- The predicates **is null** and **is not null** can be used to check for null values
- Example: find courses that do not have a webpage

```
select code  
from course  
where webpage is null;
```

- Result of expressions involving null evaluate to null
- Comparison with null returns *unknown*
- Uses same three-valued logic as relational algebra
- Aggregate functions ignore null
 - **count(*)** does *not* ignore nulls

Nested Subqueries

- A query that occurs in the **where** or **from** clause of another query is called a **subquery**
- Entire query is called **outer query** while the subquery is called **inner query** or **nested query**
- Used in tests for set membership, set cardinality, set comparisons

Set Membership

- Keyword `in` is used for set membership tests
- Example: Find faculty members who have not offered any course

Set Membership

- Keyword **in** is used for set membership tests
- Example: Find faculty members who have not offered any course

```
select *  
from faculty  
where fid not in (  
    select instructor  
    from offering );
```

Scoping of Attributes

- It is always a good practice to qualify attributes
- It is better to rename relations if it is used in both the outer and the inner queries

Scoping of Attributes

- It is always a good practice to qualify attributes
- It is better to rename relations if it is used in both the outer and the inner queries
- An unqualified attribute refers to the *innermost* query

Scoping of Attributes

- It is always a good practice to qualify attributes
- It is better to rename relations if it is used in both the outer and the inner queries
- An unqualified attribute refers to the *innermost* query
- When a nested query refers to an attribute in the outer query, it is called a **correlated query**
- Example: Find names of all students who have taken course with an instructor with the same name

Scoping of Attributes

- It is always a good practice to qualify attributes
- It is better to rename relations if it is used in both the outer and the inner queries
- An unqualified attribute refers to the *innermost* query
- When a nested query refers to an attribute in the outer query, it is called a **correlated query**
- Example: Find names of all students who have taken course with an instructor with the same name

```
select student.name
from registration as R, student as S, faculty as F
where S.roll = R.roll and S.name = F.name and F.fid in (
    select instructor
    from offering
    where offering.coursecode = R.coursecode );
```

- Inner query is evaluated for *each tuple* in the outer query

Correlated Queries

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

student	
roll	name
11	AB
12	CD
13	EF

faculty	
fid	name
101	AB
102	EF
103	GH

registration	
coursecode	roll
1	11
2	12
3	13

Correlated Queries

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

student		faculty		registration	
roll	name	fid	name	coursecode	roll
11	AB	101	AB	1	11
12	CD	102	EF	2	12
13	EF	103	GH	3	13

(R.roll = S.roll and S.name = F.name)(R × S × F)					
R.coursecode	R.roll	S.roll	S.name	F.fid	F.name
1	11	11	AB	101	AB
3	13	13	EF	102	EF

Evaluation Per Tuple

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

R.roll = S.roll and S.name = F.name			
R.coursecode	roll	name	F.fid
1	11	AB	101
3	13	EF	102

offering	
O.coursecode	instructor
1	101
2	102
3	103

- $\langle 1, 11, AB, 101 \rangle$

Evaluation Per Tuple

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

R.roll = S.roll and S.name = F.name				offering	
R.coursecode	roll	name	F.fid	O.coursecode	instructor
1	11	AB	101	1	101
3	13	EF	102	2	102
				3	103

- $\langle 1, 11, AB, 101 \rangle$
 - With R.coursecode = 1, offering chooses only instructor {101}

Evaluation Per Tuple

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

R.roll = S.roll and S.name = F.name				offering	
R.coursecode	roll	name	F.fid	O.coursecode	instructor
1	11	AB	101	1	101
3	13	EF	102	2	102
				3	103

- $\langle 1, 11, AB, 101 \rangle$
 - With $R.coursecode = 1$, offering chooses only instructor $\{101\}$
 - Now, since $fid = 101 \in \{101\}$, tuple is returned

Evaluation Per Tuple

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

R.roll = S.roll and S.name = F.name				offering	
R.coursecode	roll	name	F.fid	O.coursecode	instructor
1	11	AB	101	1	101
3	13	EF	102	2	102
				3	103

- $\langle 1, 11, AB, 101 \rangle$
 - With $R.coursecode = 1$, offering chooses only instructor $\{101\}$
 - Now, since $fid = 101 \in \{101\}$, tuple is returned
- $\langle 3, 13, EF, 102 \rangle$

Evaluation Per Tuple

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

R.roll = S.roll and S.name = F.name				offering	
R.coursecode	roll	name	F.fid	O.coursecode	instructor
1	11	AB	101	1	101
3	13	EF	102	2	102
				3	103

- $\langle 1, 11, AB, 101 \rangle$
 - With R.coursecode = 1, offering chooses only instructor {101}
 - Now, since $fid = 101 \in \{101\}$, tuple is returned
- $\langle 3, 13, EF, 102 \rangle$
 - With R.coursecode = 3, offering chooses only instructor {103}

Evaluation Per Tuple

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O  
    where O.coursecode = R.coursecode );
```

R.roll = S.roll and S.name = F.name				offering	
R.coursecode	roll	name	F.fid	O.coursecode	instructor
1	11	AB	101	1	101
3	13	EF	102	2	102
				3	103

- $\langle 1, 11, AB, 101 \rangle$
 - With R.coursecode = 1, offering chooses only instructor {101}
 - Now, since $fid = 101 \in \{101\}$, tuple is returned
- $\langle 3, 13, EF, 102 \rangle$
 - With R.coursecode = 3, offering chooses only instructor {103}
 - Now, since $fid = 102 \notin \{103\}$, tuple is not returned

Non-Correlated Query

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O, registration as G  
    where O.coursecode = G.coursecode );
```

R.roll = S.roll and S.name = F.name				Inner query
R.coursecode	roll	name	F.fid	O.instructor
1	11	AB	101	101
3	13	EF	102	102
				103

- $\langle 1, 11, AB, 101 \rangle$

Non-Correlated Query

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O, registration as G  
    where O.coursecode = G.coursecode );
```

R.roll = S.roll and S.name = F.name				Inner query
R.coursecode	roll	name	F.fid	O.instructor
1	11	AB	101	101
3	13	EF	102	102
				103

- $\langle 1, 11, AB, 101 \rangle$
 - Since $fid = 101 \in \{101, 102, 103\}$, tuple is returned

Non-Correlated Query

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O, registration as G  
    where O.coursecode = G.coursecode );
```

R.roll = S.roll and S.name = F.name				Inner query
R.coursecode	roll	name	F.fid	O.instructor
1	11	AB	101	101
3	13	EF	102	102
				103

- $\langle 1, 11, AB, 101 \rangle$
 - Since $fid = 101 \in \{101, 102, 103\}$, tuple is returned
- $\langle 3, 13, EF, 102 \rangle$

Non-Correlated Query

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O, registration as G  
    where O.coursecode = G.coursecode );
```

R.roll = S.roll and S.name = F.name				Inner query
R.coursecode	roll	name	F.fid	O.instructor
1	11	AB	101	101
3	13	EF	102	102
				103

- $\langle 1, 11, AB, 101 \rangle$
 - Since $fid = 101 \in \{101, 102, 103\}$, tuple is returned
- $\langle 3, 13, EF, 102 \rangle$
 - Since $fid = 102 \in \{101, 102, 103\}$, tuple is returned

Non-Correlated Query

```
select S.name  
from registration as R, student as S, faculty as F  
where S.roll = R.roll and S.name = F.name and F.fid in (  
    select O.instructor  
    from offering as O, registration as G  
    where O.coursecode = G.coursecode );
```

R.roll = S.roll and S.name = F.name				Inner query
R.coursecode	roll	name	F.fid	O.instructor
1	11	AB	101	101
3	13	EF	102	102
				103

- $\langle 1, 11, AB, 101 \rangle$
 - Since $fid = 101 \in \{101, 102, 103\}$, tuple is returned
- $\langle 3, 13, EF, 102 \rangle$
 - Since $fid = 102 \in \{101, 102, 103\}$, tuple is returned
- Thus, non-correlated query results in an error

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} =$

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} =$

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} = \text{false}$
 - $5 = \text{some}\{0, 5\} =$

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} = \text{false}$
 - $5 = \text{some}\{0, 5\} = \text{true}$
 - $5 \neq \text{some}\{0, 5\} =$

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} = \text{false}$
 - $5 = \text{some}\{0, 5\} = \text{true}$
 - $5 \neq \text{some}\{0, 5\} = \text{true}$

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} = \text{false}$
 - $5 = \text{some}\{0, 5\} = \text{true}$
 - $5 \neq \text{some}\{0, 5\} = \text{true}$
- $(= \text{ some}) \equiv (\text{in})$
- $(\neq \text{ some}) \equiv (\text{not in})$

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} = \text{false}$
 - $5 = \text{some}\{0, 5\} = \text{true}$
 - $5 \neq \text{some}\{0, 5\} = \text{true}$
- $(= \text{ some}) \equiv (\text{in})$
- $(\neq \text{ some}) \equiv (\text{not in})$
- Example: Find roll numbers of students who have CPI greater than some student in “CSE”

Set Comparison: some

- $(F\langle\text{comp}\rangle \text{ some } r) \Leftrightarrow (\exists t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{some}\{0, 5, 6\} = \text{true}$
 - $5 < \text{some}\{0, 5\} = \text{false}$
 - $5 = \text{some}\{0, 5\} = \text{true}$
 - $5 \neq \text{some}\{0, 5\} = \text{true}$
- $(= \text{some}) \equiv (\text{in})$
- $(\neq \text{some}) \equiv (\text{not in})$
- Example: Find roll numbers of students who have CPI greater than some student in “CSE”

```
select roll
from student
where cpi > some (
    select cpi
    from student
    where dept = 'CSE' );
```

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} =$

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} =$

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} = \text{true}$
 - $5 = \text{all}\{0, 5\} =$

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} = \text{true}$
 - $5 = \text{all}\{0, 5\} = \text{false}$
 - $5 \neq \text{all}\{4, 6\} =$

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} = \text{true}$
 - $5 = \text{all}\{0, 5\} = \text{false}$
 - $5 \neq \text{all}\{4, 6\} = \text{true}$

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} = \text{true}$
 - $5 = \text{all}\{0, 5\} = \text{false}$
 - $5 \neq \text{all}\{4, 6\} = \text{true}$
- $(\neq \text{ all}) \equiv (\text{not in})$
- $(= \text{ all}) \neq (\text{in})$

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} = \text{true}$
 - $5 = \text{all}\{0, 5\} = \text{false}$
 - $5 \neq \text{all}\{4, 6\} = \text{true}$
- $(\neq \text{ all}) \equiv (\text{not in})$
- $(= \text{ all}) \neq (\text{in})$
- Example: Find roll numbers of students who have CPI greater than all students in “CSE”

Set Comparison: all

- $(F\langle\text{comp}\rangle \text{ all } r) \Leftrightarrow (\forall t \in r (F\langle\text{comp}\rangle t))$
- Examples:
 - $5 < \text{all}\{0, 5, 6\} = \text{false}$
 - $5 < \text{all}\{6, 9\} = \text{true}$
 - $5 = \text{all}\{0, 5\} = \text{false}$
 - $5 \neq \text{all}\{4, 6\} = \text{true}$
- $(\neq \text{ all}) \equiv (\text{not in})$
- $(= \text{ all}) \neq (\text{in})$
- Example: Find roll numbers of students who have CPI greater than all students in “CSE”

```
select roll
from student
where cpi > all (
    select cpi
    from student
    where dept = 'CSE' );
```

Empty Set

- **exists** tests if the relation is empty
- $(\text{exists } r) \Leftrightarrow (r \neq \Phi)$
- $(\text{not exists } r) \Leftrightarrow (r = \Phi)$

Empty Set

- **exists** tests if the relation is empty
- $(\text{exists } r) \Leftrightarrow (r \neq \Phi)$
- $(\text{not exists } r) \Leftrightarrow (r = \Phi)$
- Example: Find faculty members who have offered courses in 2018

Empty Set

- **exists** tests if the relation is empty
- $(\text{exists } r) \Leftrightarrow (r \neq \Phi)$
- $(\text{not exists } r) \Leftrightarrow (r = \Phi)$
- Example: Find faculty members who have offered courses in 2018

```
select fid
from faculty F
where exists (
    select instructor
    from offering O
    where yr = 2018 and O.instructor = F.fid );
```

Empty Set

- **exists** tests if the relation is empty
- $(\text{exists } r) \Leftrightarrow (r \neq \Phi)$
- $(\text{not exists } r) \Leftrightarrow (r = \Phi)$
- Example: Find faculty members who have offered courses in 2018

```
select fid
from faculty F
where exists (
    select instructor
    from offering O
    where yr = 2018 and O.instructor = F.fid );
```

- Example: Find faculty members who have not offered courses in 2018

Empty Set

- **exists** tests if the relation is empty
- $(\text{exists } r) \Leftrightarrow (r \neq \Phi)$
- $(\text{not exists } r) \Leftrightarrow (r = \Phi)$
- Example: Find faculty members who have offered courses in 2018

```
select fid
from faculty F
where exists (
    select instructor
    from offering O
    where yr = 2018 and O.instructor = F.fid );
```

- Example: Find faculty members who have not offered courses in 2018

```
...
where not exists (
    ... );
```

Duplication in Sets

- **unique** tests if the relation contains duplicate tuples
- $(\text{unique } r) \Leftrightarrow (\forall t, s \in r (t \neq s))$
- $(\text{not unique } r) \Leftrightarrow (\exists t, s \in r (t = s))$

Duplication in Sets

- **unique** tests if the relation contains duplicate tuples
- $(\text{unique } r) \Leftrightarrow (\forall t, s \in r (t \neq s))$
- $(\text{not unique } r) \Leftrightarrow (\exists t, s \in r (t = s))$
- Example: Find faculty members who have only offered one course

Duplication in Sets

- **unique** tests if the relation contains duplicate tuples
- $(\text{unique } r) \Leftrightarrow (\forall t, s \in r (t \neq s))$
- $(\text{not unique } r) \Leftrightarrow (\exists t, s \in r (t = s))$
- Example: Find faculty members who have only offered one course

```
select fid
from faculty as F
where unique (
    select coursecode
    from offering O
    where O.instructor = F.fid );
```

- Example: Find faculty members who have offered multiple courses

Duplication in Sets

- **unique** tests if the relation contains duplicate tuples
- $(\text{unique } r) \Leftrightarrow (\forall t, s \in r (t \neq s))$
- $(\text{not unique } r) \Leftrightarrow (\exists t, s \in r (t = s))$
- Example: Find faculty members who have only offered one course

```
select fid
from faculty as F
where unique (
    select coursecode
    from offering O
    where O.instructor = F.fid );
```

- Example: Find faculty members who have offered multiple courses

```
...
where not unique (
    ... );
```

Explicit Sets

- Use set literals specified within brackets
- Example: Find students in “CSE” and “ECO” 33

```
select *  
from student  
where dept in ( ' 'CSE' ' , ' 'ECO' ' );
```

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order
 - 1 **from**

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order
 - 1 **from**
 - 2 **where**

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order
 - 1 **from**
 - 2 **where**
 - 3 **group by**

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order
 - 1 **from**
 - 2 **where**
 - 3 **group by**
 - 4 **having**

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order
 - 1 **from**
 - 2 **where**
 - 3 **group by**
 - 4 **having**
 - 5 **select**

Summary of SQL Query Format

- May contain up to six clauses
- May be nested
- Only the first two, **select** and **from**, are mandatory
- Format (in order)
 - select** < attribute list >
 - from** < relation list >
 - where** < predicate or tuple condition >
 - group by** < group attribute list >
 - having** < group condition >
 - order by** < attribute list >
- Execution order
 - 1 **from**
 - 2 **where**
 - 3 **group by**
 - 4 **having**
 - 5 **select**
 - 6 **order by**

Derived Relations

- In the **from** clause, a derived relation (result of a subquery) can be used

Derived Relations

- In the **from** clause, a derived relation (result of a subquery) can be used
- Example: Find departments and average CPIs where average CPI is greater than 8.0

Derived Relations

- In the **from** clause, a derived relation (result of a subquery) can be used
- Example: Find departments and average CPIs where average CPI is greater than 8.0

```
select deptid , avg_cpi
from (
    select dept , avg(cpi)
    from student
    group by dept )
as dept_avg (deptid , avg_cpi)
where avg_cpi >= 8.0;
```

Derived Relations

- In the **from** clause, a derived relation (result of a subquery) can be used
- Example: Find departments and average CPIs where average CPI is greater than 8.0

```
select deptid , avg_cpi
from (
    select dept , avg(cpi)
    from student
    group by dept )
as dept_avg (deptid , avg_cpi)
where avg_cpi >= 8.0;
```

- Avoids using **having** clause

With

- **with** clause defines a temporary relation
- This temporary relation is available *only* to the query using the **with** clause

With

- **with** clause defines a temporary relation
- This temporary relation is available *only* to the query using the **with** clause
- Example: Find departments and average CPIs where average CPI is greater than 8.0

With

- **with** clause defines a temporary relation
- This temporary relation is available *only* to the query using the **with** clause
- Example: Find departments and average CPIs where average CPI is greater than 8.0

```
with dept_avg (deptid, avg_cpi) as  
    select dept, avg(cpi)  
    from student  
    group by dept  
select deptid, avg_cpi  
from dept_avg  
where avg_cpi >= 8.0;
```

Insertion

Insertion

- `insert into ... values` statement

Insertion

- `insert into ... values` statement
- Example: Create a new student “ABC” with roll 1897 and department 7

Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- May omit schema

```
insert into student  
values (1897, ‘‘ABC’’, 7, 0.0);
```

Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- May omit schema

```
insert into student  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- If value is not known, specify *null*

```
insert into student  
values (1897, ‘‘ABC’’, 7, null);
```

Insertion

- **insert into ... values** statement
- Example: Create a new student “ABC” with roll 1897 and department 7

```
insert into student(roll , name, dept, cpi)  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- May omit schema

```
insert into student  
values (1897, ‘‘ABC’’, 7, 0.0);
```

- If value is not known, specify *null*

```
insert into student  
values (1897, ‘‘ABC’’, 7, null);
```

- To avoid *null*, specify schema

```
insert into student(roll , name, dept)  
values (1897, ‘‘ABC’’, 7);
```

Insertion (contd.)

- Example: Create a course of code 9 for every department with the same type

Insertion (contd.)

- Example: Create a course of code 9 for every department with the same type

```
insert into course(code, title , webpage, ctype)
select 9, 'New', null, type
from course
where type in (
    select deptid
    from department );
```

- Query is evaluated fully before any tuple is inserted

Insertion (contd.)

- Example: Create a course of code 9 for every department with the same type

```
insert into course(code, title , webpage, ctype)
select 9, 'New', null, type
from course
where type in (
    select deptid
    from department );
```

- Query is evaluated fully before any tuple is inserted
- Otherwise, infinite insertion happens for queries like

```
insert into r
select * from r;
```

Deletion

- `delete from ... where` statement

Deletion

- `delete from ... where` statement
- Example: Delete student with roll number 1946

Deletion

- **delete from ... where** statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```
- **where** selects tuples that will be deleted

Deletion

- `delete from ... where` statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```

- `where` selects tuples that will be deleted
- If `where` is empty,

Deletion

- **delete from ... where** statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```

- **where** selects tuples that will be deleted
- If **where** is empty, all tuples are deleted

Deletion

- **delete from ... where** statement
- Example: Delete student with roll number 1946

```
delete from student  
where roll = 1946;
```

- **where** selects tuples that will be deleted
- If **where** is empty, all tuples are deleted
- Delete all students

```
delete from student;
```

Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

```
delete from student  
where cpi < (  
    select avg(cpi)  
    from student );
```

Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

```
delete from student  
where cpi < (  
    select avg(cpi)  
    from student );
```

- Average is computed before any tuple is deleted
- It is *not* re-computed

Deletion (contd.)

- Example: Delete all students whose CPI is less than the average CPI

```
delete from student  
where cpi < (  
    select avg(cpi)  
    from student );
```

- Average is computed before any tuple is deleted
- It is *not* re-computed
- Otherwise, average keeps changing
- Ultimately, only the student with the largest CPI remains

Updating

- `update ... set ... where` statement
- `where` selects tuples that will be updated

Updating

- `update ... set ... where` statement
- `where` selects tuples that will be updated
- Example: Update value of grade 'E' to 2

Updating

- **update ... set ... where** statement
- **where** selects tuples that will be updated
- Example: Update value of grade 'E' to 2

```
update grade  
set value = 2.0  
where gradecode = 'E';
```

Updating

- **update ... set ... where** statement
- **where** selects tuples that will be updated
- Example: Update value of grade 'E' to 2

```
update grade  
set value = 2.0  
where gradecode = 'E';
```

- If **where** is empty, all tuples are updated with the new value
- Example: Increase CPI of all students by 5%

Updating

- **update ... set ... where** statement
- **where** selects tuples that will be updated
- Example: Update value of grade 'E' to 2

```
update grade  
set value = 2.0  
where gradecode = 'E';
```

- If **where** is empty, all tuples are updated with the new value
- Example: Increase CPI of all students by 5%

```
update student  
set cpi = cpi * 1.05;
```

Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important

Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important
- **case** statement handles conditional updates in a better manner and is sometimes necessary

Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;
```

```
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important
- **case** statement handles conditional updates in a better manner and is sometimes necessary
- Example: Increase CPI of all students by 10% where CPI is less than 6.0, by 5% when less than 8.0, and 2% otherwise

Updating (contd.)

- Example: Increase CPI of all students by 10% where CPI is less than 6.0 and by 5% otherwise

```
update student set cpi = cpi * 1.05 where cpi >= 6.0;  
update student set cpi = cpi * 1.10 where cpi < 6.0;
```

- Order of statements is important
- **case** statement handles conditional updates in a better manner and is sometimes necessary
- Example: Increase CPI of all students by 10% where CPI is less than 6.0, by 5% when less than 8.0, and 2% otherwise

```
update student  
set cpi =  
  case (cpi)  
    when cpi < 6.0 then cpi * 1.10  
    when cpi < 8.0 then cpi * 1.05  
    else cpi * 1.02  
end;
```

Join

- Join types: **inner join**, **left (outer) join**, **right (outer) join**, **full (outer) join**
- Join conditions: **natural**, **on** \langle predicate \rangle , **using** (\langle attribute list \rangle)
- Examples

```
student inner join program on student.roll = program.roll;
```

Join

- Join types: **inner join**, **left (outer) join**, **right (outer) join**, **full (outer) join**
- Join conditions: **natural**, **on** \langle predicate \rangle , **using** (\langle attribute list \rangle)
- Examples

```
student inner join program on student.roll = program.roll;  
student natural left join program;
```

Join

- Join types: **inner join**, **left (outer) join**, **right (outer) join**, **full (outer) join**
- Join conditions: **natural**, **on** < predicate >, **using** (< attribute list >)
- Examples

```
student inner join program on student.roll = program.roll;  
student natural left join program;  
student right outer join program using (roll);
```

Join

- Join types: **inner join**, **left (outer) join**, **right (outer) join**, **full (outer) join**
- Join conditions: **natural**, **on** \langle predicate \rangle , **using** (\langle attribute list \rangle)
- Examples

```
student inner join program on student.roll = program.roll;  
student natural left join program;  
student right outer join program using (roll);
```
- Inner natural join is assumed when nothing is mentioned

Join

- Join types: **inner join**, **left (outer) join**, **right (outer) join**, **full (outer) join**
- Join conditions: **natural**, **on** < predicate >, **using** (< attribute list >)
- Examples

```
student inner join program on student.roll = program.roll;  
student natural left join program;  
student right outer join program using (roll);
```

- Inner natural join is assumed when nothing is mentioned
- Multiple relations can be joined

```
student join program join registration;
```


Views

- A relation that is not present physically but is made visible to the user is called a **view**
- A view is a *virtual* relation derived from other relations

Views

- A relation that is not present physically but is made visible to the user is called a **view**
- A view is a *virtual* relation derived from other relations
- It helps in query processing
 - If a sub-query is very common, obtain a view for it
- It helps in hiding certain data from a user
 - A view can leave out sensitive attributes
 - Example:

```
create view student_program as  
student natural join program;
```

Views

- A relation that is not present physically but is made visible to the user is called a **view**
- A view is a *virtual* relation derived from other relations
- It helps in query processing
 - If a sub-query is very common, obtain a view for it
- It helps in hiding certain data from a user
 - A view can leave out sensitive attributes
 - Example:

```
create view student_program as  
student natural join program;
```

- A view can be deleted simply using **drop**

```
drop student_program;
```
- A view has full query capabilities, but limited modification facilities
- A view can be defined using other views, but not itself

Storing Views

- A view is *not* stored physically
- Only the query expression is stored
- Wherever a view is used, the query expression is substituted

Storing Views

- A view is *not* stored physically
- Only the query expression is stored
- Wherever a view is used, the query expression is substituted
- Example: Find students and corresponding programs for “CSE”

Storing Views

- A view is *not* stored physically
- Only the query expression is stored
- Wherever a view is used, the query expression is substituted
- Example: Find students and corresponding programs for “CSE”

```
select *  
from student_program  
where dept = ‘ ‘CSE’ ’;
```

is expanded at runtime to

Storing Views

- A view is *not* stored physically
- Only the query expression is stored
- Wherever a view is used, the query expression is substituted
- Example: Find students and corresponding programs for “CSE”

```
select *  
from student_program  
where dept = ‘‘CSE’’;
```

is expanded at runtime to

```
select *  
from  
    ( student natural join program )  
where dept = ‘‘CSE’’;
```

Storing Views

- A view is *not* stored physically
- Only the query expression is stored
- Wherever a view is used, the query expression is substituted
- Example: Find students and corresponding programs for “CSE”

```
select *  
from student_program  
where dept = ‘‘CSE’’;
```

is expanded at runtime to

```
select *  
from  
    ( student natural join program )  
where dept = ‘‘CSE’’;
```

- This allows to capture all updates in the base relations

Storing Views

- A view is *not* stored physically
- Only the query expression is stored
- Wherever a view is used, the query expression is substituted
- Example: Find students and corresponding programs for “CSE”

```
select *  
from student_program  
where dept = ‘‘CSE’’;
```

is expanded at runtime to

```
select *  
from  
    ( student natural join program )  
where dept = ‘‘CSE’’;
```

- This allows to capture all updates in the base relations
- If a view is **materialized**, it is stored physically
- To ensure consistency, database *must* update materialized views once base relations are updated

Updating a View

- Updating a view causes many problems, and is, in general, not allowed

Updating a View

- Updating a view causes many problems, and is, in general, not allowed
- Update must map to updates on the base relations

Updating a View

- Updating a view causes many problems, and is, in general, not allowed
- Update must map to updates on the base relations
- If a view involves join or Cartesian product, update must map to updates on all the base relations

Updating a View

- Updating a view causes many problems, and is, in general, not allowed
- Update must map to updates on the base relations
- If a view involves join or Cartesian product, update must map to updates on all the base relations
 - Not always possible

Updating a View

- Updating a view causes many problems, and is, in general, not allowed
- Update must map to updates on the base relations
- If a view involves join or Cartesian product, update must map to updates on all the base relations
 - Not always possible
- Problems with **insert** or **delete**

Updating a View

- Updating a view causes many problems, and is, in general, not allowed
- Update must map to updates on the base relations
- If a view involves join or Cartesian product, update must map to updates on all the base relations
 - Not always possible
- Problems with **insert** or **delete**
 - Spurious tuple
 - Null
 - Non-uniqueness

Triggers

- A **trigger** statement allows *automatic* and *active* management during database modifications
- It is invoked *only* by the database engine and not by the user

Triggers

- A **trigger** statement allows *automatic* and *active* management during database modifications
- It is invoked *only* by the database engine and not by the user
- It follows the **event-condition-action (ECA)** model
 - **Event**: Database modification
 - **Condition**: Invoked only if true; if no condition, then assumed true
 - **Action**: Database action or any program

Triggers

- A **trigger** statement allows *automatic* and *active* management during database modifications
- It is invoked *only* by the database engine and not by the user
- It follows the **event-condition-action (ECA)** model
 - **Event**: Database modification
 - **Condition**: Invoked only if true; if no condition, then assumed true
 - **Action**: Database action or any program
- It may not allow the full range of modification statements

Triggers

- A **trigger** statement allows *automatic* and *active* management during database modifications
- It is invoked *only* by the database engine and not by the user
- It follows the **event-condition-action (ECA)** model
 - **Event**: Database modification
 - **Condition**: Invoked only if true; if no condition, then assumed true
 - **Action**: Database action or any program
- It may not allow the full range of modification statements
- It can be called *before* or *after* the modification
- New and old values are referenced using **new** and **old** keywords
 - **new** refers to a inserted or new value of updated tuple
 - **old** refers to a deleted or old value of updated tuple
- By default, it is for each row (i.e., tuple)

Creating Triggers

- Created using a `create trigger` command

Creating Triggers

- Created using a `create trigger` command
- Example: Update the coursecode of offering when the code of a course of type 9 is updated

Creating Triggers

- Created using a `create trigger` command
- Example: Update the coursecode of offering when the code of a course of type 9 is updated

```
create trigger update_code
after update of code on course
for each row
when ctype = 9
begin
    update offering set coursecode = new.code where
        coursecode = old.code;
end;
```

Creating Triggers

- Created using a **create trigger** command
- Example: Update the coursecode of offering when the code of a course of type 9 is updated

```
create trigger update_code
  after update of code on course
  for each row
  when ctype = 9
  begin
    update offering set coursecode = new.code where
      coursecode = old.code;
  end;
```

- A trigger can be deleted simply using **drop**
- ```
drop update_code ;
```

# Indices in SQL

- **create [unique] index name on r (a, b)** creates an index *name* on attributes *a*, *b* of relation *r*
  - *unique* does not allow duplicates on the attributes

**create index** idx\_ctype **on** course (ctype);



# Indices in SQL

- `create [unique] index name on r (a, b)` creates an index *name* on attributes *a*, *b* of relation *r*
  - *unique* does not allow duplicates on the attributes

**create index** idx\_ctype **on** course (ctype);

- Index will be used whenever attribute *a* of relation *r* is used
- Often, primary keys are implicitly indexed

# Indices in SQL

- `create [unique] index name on r (a, b)` creates an index *name* on attributes *a*, *b* of relation *r*
  - *unique* does not allow duplicates on the attributes

**create index** idx\_cotype **on** course (cotype);

- Index will be used whenever attribute *a* of relation *r* is used
- Often, primary keys are implicitly indexed
- Slows down modification operations as index is also modified

# Indices in SQL

- **create [unique] index name on r (a, b)** creates an index *name* on attributes *a*, *b* of relation *r*
  - *unique* does not allow duplicates on the attributes

```
create index idx_cotype on course (cotype);
```

- Index will be used whenever attribute *a* of relation *r* is used
- Often, primary keys are implicitly indexed
- Slows down modification operations as index is also modified
- **drop index i** deletes the index

```
drop index idx_cotype;
```

# Access Control

- SQL is a **data control language (DCL)**
- Can be used to control accesses of users to data
- *Data security*

# Access Control

- SQL is a **data control language (DCL)**
- Can be used to control accesses of users to data
- *Data security*
- **grant**: **grant privilege on object to user [with grant option]**
  - *privilege*: Can be **all** or specific:
    - System privileges on tables and views: **create**, **alter**, **drop**
    - Object privileges: **select**, **insert/update/delete**, **execute**
  - *object*: table, view, stored procedure
  - *user*: **public** or particular username or *role*
  - *with grant option*: can grant access rights to others

# Access Control

- SQL is a **data control language (DCL)**
- Can be used to control accesses of users to data
- *Data security*
- **grant**: **grant privilege on object to user [with grant option]**
  - *privilege*: Can be **all** or specific:
    - System privileges on tables and views: **create**, **alter**, **drop**
    - Object privileges: **select**, **insert/update/delete**, **execute**
  - *object*: table, view, stored procedure
  - *user*: **public** or particular username or *role*
  - *with grant option*: can grant access rights to others
- **revoke**: **revoke privilege on object from user**

# Access Control

- SQL is a **data control language (DCL)**
- Can be used to control accesses of users to data
- *Data security*
- **grant**: **grant privilege on object to user [with grant option]**
  - *privilege*: Can be **all** or specific:
    - System privileges on tables and views: **create**, **alter**, **drop**
    - Object privileges: **select**, **insert/update/delete**, **execute**
  - *object*: table, view, stored procedure
  - *user*: **public** or particular username or *role*
  - *with grant option*: can grant access rights to others
- **revoke**: **revoke privilege on object from user**

**grant select on student to xyz;**

**grant all on course to abc with grant option;**

**revoke create on student from xyz;**

# Transaction Control

- **Transactions** are groups of statements that are executed atomically



# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction

# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction
- After modification operations, a transaction can `commit` or `rollback`

# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction
- After modification operations, a transaction can `commit` or `rollback`
- Within a transaction, a checkpoint can be set by `savepoint savepoint-name`

# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction
- After modification operations, a transaction can `commit` or `rollback`
- Within a transaction, a checkpoint can be set by `savepoint savepoint-name`
- A transaction can `rollback to` a particular `savepoint-name`

# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction
- After modification operations, a transaction can `commit` or `rollback`
- Within a transaction, a checkpoint can be set by `savepoint savepoint-name`
- A transaction can `rollback to` a particular `savepoint-name`
- A savepoint can be removed by `release savepoint savepoint-name`

# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction
- After modification operations, a transaction can `commit` or `rollback`
- Within a transaction, a checkpoint can be set by `savepoint savepoint-name`
- A transaction can `rollback to` a particular `savepoint-name`
- A savepoint can be removed by `release savepoint savepoint-name`

| roll | name |
|------|------|
| 1    | AB   |
| 2    | CD   |

```
set transaction read write ;
delete from student where roll = 1 ;
commit ;
```

# Transaction Control

- **Transactions** are groups of statements that are executed atomically
- `set transaction [read write | read only]` starts a transaction
- After modification operations, a transaction can `commit` or `rollback`
- Within a transaction, a checkpoint can be set by `savepoint savepoint-name`
- A transaction can `rollback to` a particular `savepoint-name`
- A savepoint can be removed by `release savepoint savepoint-name`

| roll | name |
|------|------|
| 1    | AB   |
| 2    | CD   |

```
set transaction read write;
delete from student where roll = 1;
commit;
```

| roll | name |
|------|------|
| 2    | CD   |

# Savepoint Example

| roll | name |
|------|------|
| 1    | AB   |
| 2    | CD   |
| 3    | EF   |
| 4    | GH   |

```
set transaction read write;
savepoint sp1;
delete from student where roll = 1;
savepoint sp2;
delete from student where roll = 2;
rollback to sp2;
```



# Savepoint Example

| roll | name |
|------|------|
| 1    | AB   |
| 2    | CD   |
| 3    | EF   |
| 4    | GH   |

```
set transaction read write;
savepoint sp1;
delete from student where roll = 1;
savepoint sp2;
delete from student where roll = 2;
rollback to sp2;
```

| roll | name |
|------|------|
| 2    | CD   |
| 3    | EF   |
| 4    | GH   |

# Variants in SQL

- SQL standards have evolved a lot over the years
- Different vendors provide different flavors and may not implement every feature